

# Data Cleaning

Our housing data is provided to us in a file called `kc_house_data.csv`. First, we bring that data into a pandas dataframe and start examining it.

```
In [1]: import pandas as pd
import numpy as np
```

```
In [2]: df = pd.read_csv('data/kc_house_data.csv')
df
```

Out[2]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	water
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	
...	...	...	...	...	...	...	...	...	...
21592	263000018	5/21/2014	360000.0	3	2.50	1530	1131	3.0	
21593	6600060120	2/23/2015	400000.0	4	2.50	2310	5813	2.0	
21594	1523300141	6/23/2014	402101.0	2	0.75	1020	1350	2.0	
21595	291310100	1/16/2015	400000.0	3	2.50	1600	2388	2.0	
21596	1523300157	10/15/2014	325000.0	2	0.75	1020	1076	2.0	

In [3]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
id                21597 non-null int64
date              21597 non-null object
price             21597 non-null float64
bedrooms         21597 non-null int64
bathrooms        21597 non-null float64
sqft_living      21597 non-null int64
sqft_lot         21597 non-null int64
floors           21597 non-null float64
waterfront       19221 non-null float64
view             21534 non-null float64
condition        21597 non-null int64
grade            21597 non-null int64
sqft_above       21597 non-null int64
sqft_basement    21597 non-null object
yr_built         21597 non-null int64
yr_renovated     17755 non-null float64
zipcode          21597 non-null int64
```

In [4]: df.isna().sum()

```
Out[4]: id                0
date                0
price               0
bedrooms            0
bathrooms           0
sqft_living         0
sqft_lot            0
floors              0
waterfront          2376
view                63
condition           0
grade               0
sqft_above          0
sqft_basement       0
yr_built            0
yr_renovated        3842
zipcode             0
lat                 0
long                0
```

```
In [5]: for col in list(df.columns)[1:]:  
        print(col)  
        print(df[col].value_counts())
```

```
date  
6/23/2014    142  
6/26/2014    131  
6/25/2014    131  
7/8/2014     127  
4/27/2015    126  
...  
5/15/2015     1  
11/2/2014     1  
1/10/2015     1  
2/15/2015     1  
3/8/2015      1  
Name: date, Length: 372, dtype: int64  
price  
350000.0     172  
450000.0     172  
550000.0     159  
500000.0     152  
425000.0     150
```

We mostly have numerical objects, but with a few anomalies.

- "sqft\_basement" should be a number and yet is a string. We will also need to replace the '?' with 0, to indicate no basement.
- "date" will need to be modified to work with a linear regression. We will break the year out from that column to be its own feature, "yr\_sold".
- NaN values can be safely turned into 0, to indicate a lack of waterfront or renovations.
- 0 values in "yr\_renovated" won't work with most values being in the 1900's-2000's, so we will replace it with a "yr\_since\_renovated" column (yr\_sold-yr\_renovated, or yr\_sold-yr\_built when yr\_renovated=0).
- similarly, we will create a "yr\_since\_built" column (yr\_sold-yr\_built).

Our categorical variables are:

- waterfront
- view
- condition
- grade
- zipcode

```
In [6]: df['yr_sold'] = df.date.map(lambda x: int(x.split('/')[1]))
df
```

Out[6]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	
...	...	...	...	...	...	...	...	...	...
21592	263000018	5/21/2014	360000.0	3	2.50	1530	1131	3.0	
21593	6600060120	2/23/2015	400000.0	4	2.50	2310	5813	2.0	
21594	1523300141	6/23/2014	402101.0	2	0.75	1020	1350	2.0	
21595	291310100	1/16/2015	400000.0	3	2.50	1600	2388	2.0	
21596	1523300157	10/15/2014	325000.0	2	0.75	1020	1076	2.0	

```
In [7]: df.replace({'sqft_basement': {'?': '0.0'}}, inplace=True)
df.sqft_basement = pd.to_numeric(df.sqft_basement)
```

```
In [8]: df.fillna(0.0, inplace=True)
```

```
In [9]: df.isna().sum()
```

```
Out[9]: id          0
date            0
price           0
bedrooms        0
bathrooms       0
sqft_living     0
sqft_lot        0
floors          0
waterfront      0
view            0
condition       0
grade           0
sqft_above      0
sqft_basement   0
yr_built        0
yr_renovated    0
zipcode         0
lat             0
long            0
sqft_basement   0
```


```
In [10]: df['yr_since_renovation'] = np.where(df['yr_renovated']==0.0, df['yr_sold']-df['yr_built'],
df['yr_since_built'] = df['yr_sold'] - df['yr_built']
df['renovated'] = df.yr_renovated.map(lambda x: 1 if x>0 else 0)
```

In [11]: `df`

Out[11]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	water
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	
...	...	...	...	...	...	...	...	...	...
21592	263000018	5/21/2014	360000.0	3	2.50	1530	1131	3.0	
21593	6600060120	2/23/2015	400000.0	4	2.50	2310	5813	2.0	
21594	1523300141	6/23/2014	402101.0	2	0.75	1020	1350	2.0	
21595	291310100	1/16/2015	400000.0	3	2.50	1600	2388	2.0	
21596	1523300157	10/15/2014	325000.0	2	0.75	1020	1076	2.0	

21597 rows × 25 columns



We have a few features we want to remove:

- id
- date
- view
- sqft\_above
- sqft\_living15
- sqft\_lot15

"sqft\_living15", "sqft\_lot15" are not about the house itself, and are likely to be highly correlated with our other features anyways. "sqft\_above" is directly the difference between "sqft\_living" and "sqft\_basement", so it is also unnecessary. The description for "view" states "Has been viewed", and yet indicates numbers 1 through 3, with some NaN values. We aren't sure how to interpret this, so we have decided to exclude it. "date" has been incorporated into our "yr\_sold" column, and so is now unnecessary. the "id" column has no bearing on our modeling, so we will remove it as well.

In [12]: `df.drop(['id', 'view', 'sqft_above', 'sqft_living15', 'sqft_lot15', 'date'], axis=`

In [13]: df

Out[13]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	condition	grade
0	221900.0	3	1.00	1180	5650	1.0	0.0	3	7
1	538000.0	3	2.25	2570	7242	2.0	0.0	3	7
2	180000.0	2	1.00	770	10000	1.0	0.0	3	6
3	604000.0	4	3.00	1960	5000	1.0	0.0	5	7
4	510000.0	3	2.00	1680	8080	1.0	0.0	3	8
...	...	...	...	...	...	...	...	...	...
21592	360000.0	3	2.50	1530	1131	3.0	0.0	3	8
21593	400000.0	4	2.50	2310	5813	2.0	0.0	3	8
21594	402101.0	2	0.75	1020	1350	2.0	0.0	3	7
21595	400000.0	3	2.50	1600	2388	2.0	0.0	3	8
21596	325000.0	2	0.75	1020	1076	2.0	0.0	3	7

In [14]: df.to\_csv('data/cleaned\_data.csv', index=False)

# Exploratory Data Analysis

In order to prepare for our modeling, we will check which features are correlated with "price", check the feature's distribution, and look for collinear features.

```
In [1]: import pandas as pd
```

```
In [2]: df = pd.read_csv('data/cleaned_data.csv')
df.head()
```

Out[2]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	condition	grade	sqft_
0	221900.0	3	1.00	1180	5650	1.0	0.0	3	7	
1	538000.0	3	2.25	2570	7242	2.0	0.0	3	7	
2	180000.0	2	1.00	770	10000	1.0	0.0	3	6	
3	604000.0	4	3.00	1960	5000	1.0	0.0	5	7	
4	510000.0	3	2.00	1680	8080	1.0	0.0	3	8	

```
In [3]: df.describe()
```

Out[3]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	wa
count	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04	21597.000000	21597
mean	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04	1.494096	0
std	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04	0.539683	0
min	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02	1.000000	0
25%	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03	1.000000	0
50%	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.500000	0
75%	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04	2.000000	0
max	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000	1

## Which Features are Corolated with the Target

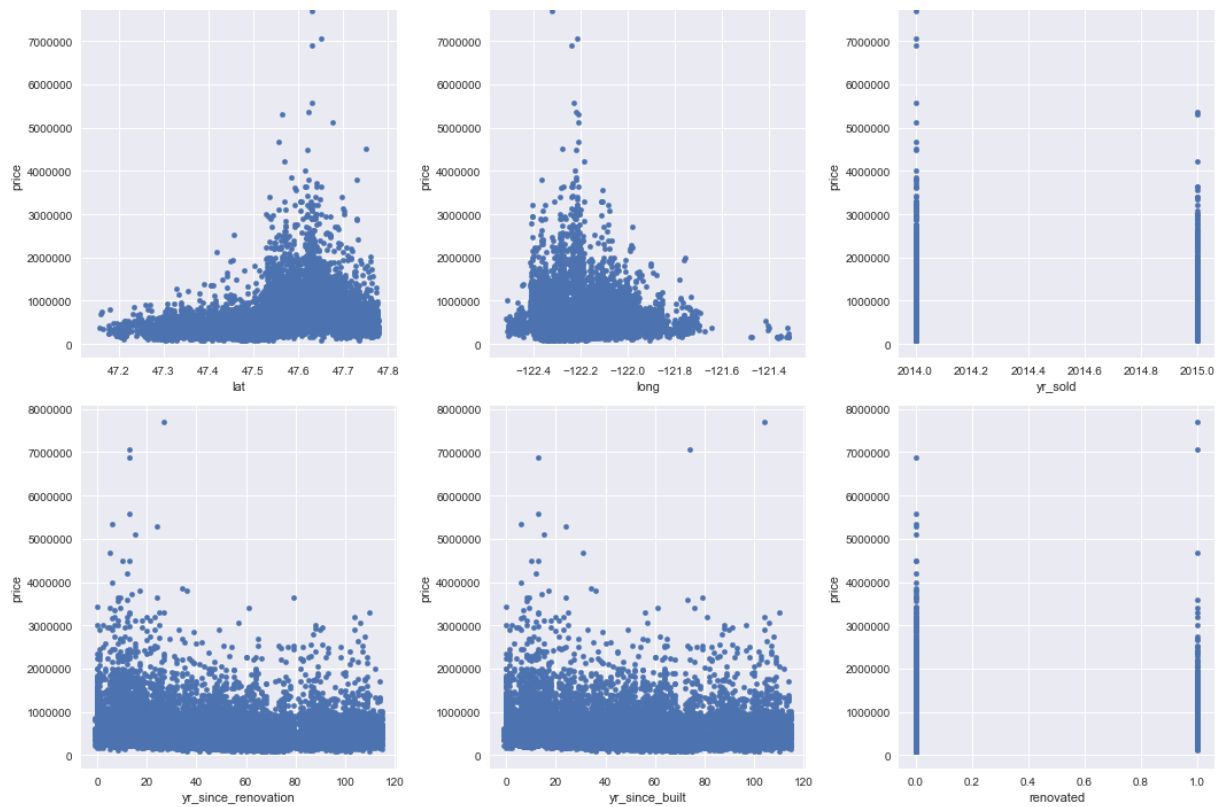
We plot each variable against "price", our target, to see what relationship they have.

```
In [4]: import matplotlib.pyplot as plt
plt.style.use('seaborn')
%matplotlib inline
```

```
In [5]: plt.figure(figsize=(15,30))
for i, col in enumerate(df.drop('price', axis=1).columns):
    ax = plt.subplot(6, 3, i+1)
    df.plot.scatter(x=col, y='price', ax=ax, legend=False)
plt.tight_layout()
plt.savefig(f'figures/scatter-plots.png')
plt.show()
```







These columns have a strong obvious correlation with price:

- bedrooms
- bathrooms
- sqft\_living
- sqft\_lot
- sqft\_basement
- yr\_built

These columns are to be treated as categorical data:

- Condition
- Grade
- Zip Code

These columns do not appear to have a strong linear correlation:

- Waterfront
- lat
- long
- yr\_sold
- yr\_since\_renovation
- yr\_since\_built
- renovated

We don't need to remove any data yet, as uncorellated data will appear with high p-values when we run our model.

# Features to Transform

A histogram of each variable will show us their distribution. If they aren't normal, they will need to be log transformed. This will also make clearer which of our variables are continuous.

```
In [6]: df.hist(figsize = (20,18))
plt.tight_layout()
plt.savefig(f'figures/histogram-plots.png')
plt.show()
```



The only seemingly normal distribution is "grade". We will need to log transform all continuous variables.

continuous variables:

- bedrooms
- bathrooms
- sqft\_living
- sqft\_lot
- sqft\_basement

- lat
- long
- yr\_since\_built

```
In [7]: df.yr_sold.value_counts()
```

```
Out[7]: 2014    14622
        2015     6975
        Name: yr_sold, dtype: int64
```

Data is only for 2 years, so this column is unlikely to help us. We will go ahead and drop this feature.

There are also two variables in particular - price and bedrooms - where the range is far greater than the average (i.e. there are many outliers). For price, this is expected, although it will need to be dealt with in order to get a more accurate model.


```
In [8]: df[df['bedrooms']>15]
```

```
Out[8]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	condition	grade	...
<b>15856</b>	640000.0	33	1.75	1620	6000	1.0	0.0	5	7	

For bedrooms, this is unexpected. It turns out this whole thing is caused by one data point. Given the square footage and number of bathrooms in the house is not near enough for 33 bedrooms, we believe this to be a typo, and will correct this to 3 bedrooms in our iterative modeling process. Here is what the bedrooms histogram should look like without this one error.

```
In [9]: df.at[15856, 'bedrooms'] = 3
        df.hist('bedrooms')
        plt.show()
```



# What Features are Colinear

Here we check which features are colinear. We don't want colinear pairs, and will drop one from each pair.


In [10]: `import seaborn as sns`

In [11]: `corr = df.drop('price' , axis=1).corr()  
  
df_corr=corr.abs().stack().reset_index().sort_values(0, ascending=False)  
df_corr['pairs'] = list(zip(df_corr.level_0, df_corr.level_1))  
df_corr.set_index(['pairs'], inplace = True)  
df_corr.drop(columns=['level_1', 'level_0'], inplace = True)  
df_corr.columns = ['cc']  
df_corr.drop_duplicates(inplace=True)  
df_corr[(df_corr.cc>.8) & (df_corr.cc <1)]`

Out[11]:

	cc
pairs	
(yr_renovated, renovated)	0.999968
(yr_built, yr_since_built)	0.999873
(yr_since_renovation, yr_since_built)	0.926424
(yr_since_renovation, yr_built)	0.926173

```
In [17]: plt.figure(figsize=(7, 7))
sns.heatmap(corr, center=0, annot=True);
plt.tight_layout()
plt.savefig(f'figures/heatmap-before.png')
plt.show()
```



In order to remove collinearity, we will drop the yr\_since\_renovation, yr\_built and the yr\_renovated columns. These are also incorporated into other features, renovated and yr\_since\_built, so they will be safe to remove.


```
In [13]: corr = df.drop(['price', 'yr_built', 'yr_renovated', 'yr_sold', 'yr_since_renovat

df_corr=corr.abs().stack().reset_index().sort_values(0, ascending=False)
df_corr['pairs'] = list(zip(df_corr.level_0, df_corr.level_1))
df_corr.set_index(['pairs'], inplace = True)
df_corr.drop(columns=['level_1', 'level_0'], inplace = True)
df_corr.columns = ['cc']
df_corr.drop_duplicates(inplace=True)
df_corr[(df_corr.cc>.8) & (df_corr.cc <1)]
```

Out[13]:

cc
<u>pairs</u>

```
In [18]: plt.figure(figsize=(7, 7))
sns.heatmap(corr, center=0, annot=True);
plt.tight_layout()
plt.savefig(f'figures/heatmap-after.png')
plt.show()
```



We check for colinearity after removing the stated features to verify we no longer have any strongly colinear features. We will be sure to remove these columns once we begin our iterative modeling process.

# Feature Engineering

Here we will perform the final steps to get our data ready for modeling, split our data into train and test portions, and create our baseline model.


In [1]: `import pandas as pd`

In [2]: `df = pd.read_csv('data/cleaned_data.csv')`  
`df`

Out[2]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	condition	grade	sqft_basement
0	221900.0	3	1.00	1180	5650	1.0	0.0	3	7	1180
1	538000.0	3	2.25	2570	7242	2.0	0.0	3	7	2570
2	180000.0	2	1.00	770	10000	1.0	0.0	3	6	770
3	604000.0	4	3.00	1960	5000	1.0	0.0	5	7	1960
4	510000.0	3	2.00	1680	8080	1.0	0.0	3	8	1680
...	...	...	...	...	...	...	...	...	...	...
21592	360000.0	3	2.50	1530	1131	3.0	0.0	3	8	1530
21593	400000.0	4	2.50	2310	5813	2.0	0.0	3	8	2310
21594	402101.0	2	0.75	1020	1350	2.0	0.0	3	7	1020
21595	400000.0	3	2.50	1600	2388	2.0	0.0	3	8	1600
21596	325000.0	2	0.75	1020	1076	2.0	0.0	3	7	1020

21597 rows × 19 columns





```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 19 columns):
price                21597 non-null float64
bedrooms             21597 non-null int64
bathrooms            21597 non-null float64
sqft_living          21597 non-null int64
sqft_lot             21597 non-null int64
floors               21597 non-null float64
waterfront           21597 non-null float64
condition            21597 non-null int64
grade                21597 non-null int64
sqft_basement        21597 non-null float64
yr_built             21597 non-null int64
yr_renovated         21597 non-null float64
zipcode              21597 non-null int64
lat                  21597 non-null float64
long                 21597 non-null float64
yr_sold              21597 non-null int64
dtype: object
```

```
In [4]: df.isna().sum()
```

```
Out[4]: price                0
bedrooms             0
bathrooms            0
sqft_living          0
sqft_lot             0
floors               0
waterfront           0
condition            0
grade                0
sqft_basement        0
yr_built             0
yr_renovated         0
zipcode              0
lat                  0
long                 0
yr_sold              0
yr_since_renovation  0
yr_since_built       0
renovated            0
dtype: int64
```

We have verified that our data has no nan values, and all data types are integers. However, we want to create dummy variables for those categories we've deemed categorical: 'floors', 'condition', 'grade', and 'zipcode'. The Pandas `get_dummies` function works on object datatypes, so we turn those columns into strings before running it.

```
In [5]: categoricals = ['floors', 'condition', 'grade', 'zipcode']
df = df.astype({col: 'str' for col in categoricals})
df = pd.get_dummies(df, drop_first=True)
df
```

Out[5]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	waterfront	sqft_basement	yr_built
0	221900.0	3	1.00	1180	5650	0.0	0.0	1955
1	538000.0	3	2.25	2570	7242	0.0	400.0	1951
2	180000.0	2	1.00	770	10000	0.0	0.0	1933
3	604000.0	4	3.00	1960	5000	0.0	910.0	1965
4	510000.0	3	2.00	1680	8080	0.0	0.0	1987
...	...	...	...	...	...	...	...	...
21592	360000.0	3	2.50	1530	1131	0.0	0.0	2009
21593	400000.0	4	2.50	2310	5813	0.0	0.0	2014
21594	402101.0	2	0.75	1020	1350	0.0	0.0	2009
21595	400000.0	3	2.50	1600	2388	0.0	0.0	2004
21596	325000.0	2	0.75	1020	1076	0.0	0.0	2008

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Columns: 103 entries, price to zipcode_98199
dtypes: float64(8), int64(7), uint8(88)
memory usage: 4.3 MB
```

Many punctuation marks won't work as column names with the statsmodel linear regression modeling, so we remove or replace them.

```
In [7]: def col_formatting(col):
        for old, new in subs:
            col = col.replace(old,new)
        return col
subs = [(' ', '_'), (',', ''), ('.', ''), ('"', ''), ('\'', ''), ('@', ''), ('+', 'plus'), (
df.columns = [col_formatting(col) for col in df.columns]
```

## Train-Test Split

We declare our train and test sets before running any modeling, using sklearn's `train_test_split` function. We keep its default of 25% of the data for the test set, and set a random state for repeatability. The data will also be saved to csv so they can be read later.

```
In [8]: from sklearn.model_selection import train_test_split
```

```
In [9]: train, test = train_test_split(df, random_state=7)
```

```
In [10]: train
```

```
Out[10]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	waterfront	sqft_basement	yr_built
<b>15200</b>	175000.0	3	1.00	1070	6164	0.0	0.0	1967
<b>20737</b>	775000.0	4	2.50	2580	5787	0.0	0.0	2007
<b>19361</b>	440000.0	4	2.50	2350	7203	0.0	0.0	1989
<b>15578</b>	1680000.0	5	5.25	4830	18707	0.0	900.0	1952
<b>8436</b>	2140000.0	4	3.75	5150	453895	0.0	790.0	1997
...	...	...	...	...	...	...	...	...
<b>919</b>	250000.0	3	2.00	1440	9220	0.0	0.0	1965
<b>20691</b>	380000.0	5	3.50	2420	4670	0.0	0.0	2013
<b>5699</b>	276500.0	4	1.75	1400	6650	0.0	0.0	1942
<b>10742</b>	440000.0	4	2.75	2340	11034	0.0	620.0	1967
<b>16921</b>	335000.0	2	1.75	1060	1202	0.0	300.0	2003

```
In [11]: test
```

```
Out[11]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	waterfront	sqft_basement	yr_built
<b>5460</b>	560000.0	4	2.75	1950	6192	0.0	0.0	1992
<b>7131</b>	500000.0	5	3.00	2920	11440	0.0	0.0	2003
<b>8759</b>	470000.0	2	1.00	1220	4000	0.0	0.0	1908
<b>14957</b>	1020000.0	4	3.00	2720	4800	0.0	930.0	1928
<b>5431</b>	375000.0	3	2.50	1930	6180	0.0	600.0	1961
...	...	...	...	...	...	...	...	...
<b>19665</b>	1850000.0	4	3.25	4160	10335	0.0	0.0	2014
<b>7714</b>	249900.0	3	1.00	1100	5000	0.0	0.0	1960
<b>2480</b>	679000.0	4	1.50	1420	4923	0.0	0.0	1928
<b>16033</b>	300000.0	3	1.75	1700	8481	0.0	0.0	1993
<b>6193</b>	450000.0	4	2.50	2070	3982	0.0	0.0	2004

```
In [12]: train.to_csv('data/train.csv', index=False)
test.to_csv('data/test.csv', index=False)
```

## Baseline Model

To create our baseline model, we will run our train set as is through the statsmodel linear regression

function. We will then collect some metrics from the model using the test data.

```
In [13]: from statsmodels.formula.api import ols
from sklearn.metrics import mean_squared_error
import numpy as np
import matplotlib.pyplot as plt
plt.style.use('seaborn')
%matplotlib inline
```

```
In [14]: predictors = '+'.join(df.columns[1:])
formula = 'price' + '~' + predictors
model = ols(formula=formula, data=train).fit()
model.summary()
```

Out[14]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.821
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.820
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	730.0
<b>Date:</b>	Sun, 29 Nov 2020	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	10:30:04	<b>Log-Likelihood:</b>	-2.1678e+05
<b>No. Observations:</b>	16197	<b>AIC:</b>	4.338e+05
<b>Df Residuals:</b>	16095	<b>BIC:</b>	4.345e+05
<b>Df Model:</b>	101		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	-9.703e+07	8.78e+06	-11.047	0.000	-1.14e+08	-7.98e+07

```
In [15]: train_r2, train_r2_adj = model.rsquared, model.rsquared_adj
train_r2, train_r2_adj
```

Out[15]: (0.8208200222124922, 0.8196956247128626)

```
In [16]: y_hat_train = model.predict(train.drop('price', axis=1))
y_train = train['price']
train_rmse = np.sqrt(mean_squared_error(y_train, y_hat_train))
train_rmse
```

Out[16]: 157156.21796974784

```
In [17]: y_hat_test = model.predict(test.drop('price', axis=1))
y_test = test['price']
test_rmse = np.sqrt(mean_squared_error(y_test, y_hat_test))
test_rmse
```

Out[17]: 139700.57671817578


```
In [18]: pvalues = model.pvalues.to_dict()
significant_items = {}
for key, value in pvalues.items():
    if value < 0.05:
        significant_items[key] = value
len(significant_items), len(pvalues)
```

Out[18]: (81, 103)

Baseline model:

- R2 of 0.821
- adjusted R2 of 0.820
- Train RMSE of 157156
- Test RMSE of 139700
- 81 significant features (p-value < 0.05)
- 103 features total

```
In [19]: import statsmodels.api as sm
import scipy.stats as stats
sm.graphics.qqplot(model.resid, dist=stats.norm, line='45', fit=True)
plt.tight_layout()
plt.savefig('figures/baseline-qq-plot.png')
plt.show()
```



Our metrics of our baseline model shows we have some problems. Our R2 isn't too bad, but with such a huge number of features, it doesn't mean much. Our RMSE values are high, and our QQ plot shows our data seems to have fat tails, although it doesn't seem to have much skew.

# Iterative Modeling Process

To start our iterative modeling process, we will create a function to print out our metrics and create our qq plot. We will have to un-log-transform before calculating RMSE in order to have a comparable metric, so that will have to be included as an argument, as well as the model itself.

```
In [1]: import pandas as pd
import statsmodels.api as sm
import scipy.stats as stats
from statsmodels.formula.api import ols
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
import numpy as np
import matplotlib.pyplot as plt
plt.style.use('seaborn-darkgrid')
%matplotlib inline
```

```
In [2]: train = pd.read_csv('data/train.csv')
test = pd.read_csv('data/test.csv')
```

```
In [3]: def get_model_data(model, is_log_transformed=False):

    train_r2, train_r2_adj = model.rsquared, model.rsquared_adj

    y_hat_train = model.predict(train.drop('price', axis=1))
    y_train = train['price']
    if is_log_transformed:
        train_rmse = np.sqrt(mean_squared_error(np.exp(y_train), np.exp(y_hat_train)))
    else:
        train_rmse = np.sqrt(mean_squared_error(y_train, y_hat_train))

    y_hat_test = model.predict(test.drop('price', axis=1))
    y_test = test['price']
    if is_log_transformed:
        test_rmse = np.sqrt(mean_squared_error(np.exp(y_test), np.exp(y_hat_test)))
    else:
        test_rmse = np.sqrt(mean_squared_error(y_test, y_hat_test))

    pvalues = model.pvalues.to_dict()
    significant_items = {}
    for key, value in pvalues.items():
        if value < 0.05:
            significant_items[key] = value

    print('R2 =', train_r2)
    print('R2 adjusted =', train_r2_adj)
    print('RMSE (train) =', train_rmse)
    print('RMSE (test) =', test_rmse)
    print('number of significant features =', len(significant_items))

    sm.graphics.qqplot(model.resid, dist=stats.norm, line='45', fit=True)
    plt.title('Q-Q Plot')
```


## Dropping Colinear Features

Here we drop the colinear features, as established in our exploratory data analysis.

```
In [4]: train.drop(['yr_built', 'yr_renovated', 'yr_sold', 'yr_since_renovation'], axis=1)
test.drop(['yr_built', 'yr_renovated', 'yr_sold', 'yr_since_renovation'], axis=1)
```

```
In [5]: predictors = '+'.join(train.columns[1:])  
formula = 'price' + '~' + predictors  
model = ols(formula=formula, data=train).fit()  
get_model_data(model)
```

```
R2 = 0.8189229383789512  
R2 adjusted = 0.8178205932404954  
RMSE (train) = 157985.98012359045  
RMSE (test) = 140609.21109533816  
number of significant features = 78
```



## Removing Outliers

Our data has a lot of high outliers in the price variable, and one very high bedroom count that is likely a typo. We will fix the one error, then remove the outliers higher than three times the standard deviation away from the mean. This will limit the range of house prices we can predict for, but will make for a much more accurate model. The mean and standard deviation are taken from our exploratory data analysis.




```
In [6]: train[train['bedrooms']>15]
```

```
Out[6]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	waterfront	sqft_basement	lat	long	yr_built	yr_renovated
4196	640000.0	33	1.75	1620	6000	0.0	580.0	47.6878	-122.3294	1965	2000

1 rows × 99 columns




```
In [7]: train.at[4196, 'bedrooms'] = 3
train[train['bedrooms']>15]
```

```
Out[7]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	waterfront	sqft_basement	lat	long	yr_built	yr_renovated
--	-------	----------	-----------	-------------	----------	------------	---------------	-----	------	----------	--------------

0 rows × 99 columns



```
In [8]: mean = 5.402966e+05
std = 3.673681e+05
lower_cutoff = mean - (3*std)
upper_cutoff = mean + (3*std)
lower_cutoff, upper_cutoff
```

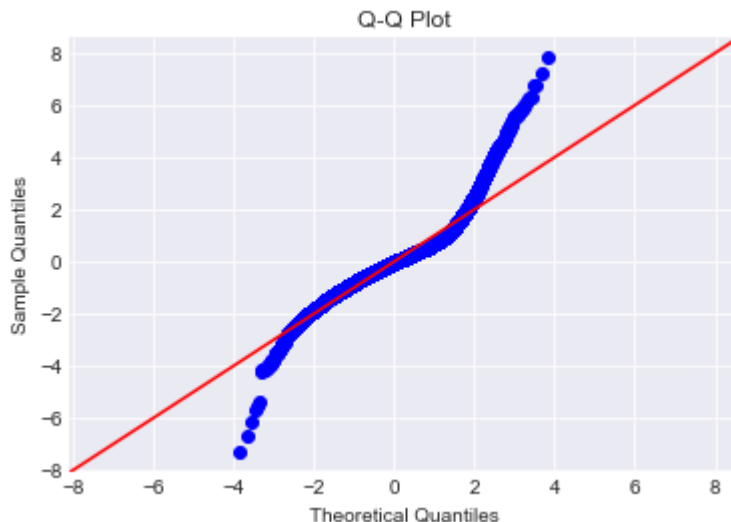
```
Out[8]: (-561807.69999999998, 1642400.9)
```

Our lower cutoff is negative. Since we have no negative prices (our minimum from our exploratory data analysis was \$78,000, and negative prices just don't make sense), we don't have to worry about our lower cutoff.

```
In [9]: train = train[train['price'] < upper_cutoff]
test = test[test['price'] < upper_cutoff]
```

```
In [10]: predictors = '+'.join(train.columns[1:])  
formula = 'price' + '~' + predictors  
model = ols(formula=formula, data=train).fit()  
get_model_data(model)
```

```
R2 = 0.8264414496035375  
R2 adjusted = 0.8253755271400415  
RMSE (train) = 107801.0153531246  
RMSE (test) = 105098.35236656363  
number of significant features = 82
```



The model looks a lot better after dropping values higher than three times the standard deviation. However, now we only have data within three times the standard deviation to work with. Our model will now only accurately predict house prices up to \$1,642,400.

## Log Transforming Continuous Variables

We will log transform most of our continuous variables in order to make them more normally distributed. We won't transform longitude and latitude, as they can have negative values. `yr_since_built` and `sqft_basement`, as well, can be zero. Log transformations do not work on negative and zero values, so we will skip these columns.

The price column, our dependent variable, will be log transformed as well. This will get us better results, but it means we need to do the inverse transformation when calculating rmse values in order to get a comparable result.



## Removing Insignificant Features

Here we loop through our model multiple times. Each time we find the feature with the highest p-value and remove it if it is higher than 0.05. By doing this, we get rid of 9 insignificant features. By reviewing the model summary, we can verify that we are left with only significant ones.

```
In [14]: model_dict = list(dict(model.pvalues).items())
model_dict.sort(key = lambda x: x[1], reverse=True)
highest_pvalue = model_dict[0]

while highest_pvalue[1] > 0.05:
    print(f'Dropping "{highest_pvalue[0]}" with p-value {highest_pvalue[1]}')
    train.drop(highest_pvalue[0], inplace = True, axis = 1)
    test.drop(highest_pvalue[0], inplace = True, axis = 1)


    predictors = '+'.join(train.columns[1:])
    formula = 'price' + '~' + predictors
    model = ols(formula=formula, data=train).fit()

    model_dict = list(dict(model.pvalues).items())
    model_dict.sort(key = lambda x: x[1], reverse=True)
    highest_pvalue = model_dict[0]
```

```
Dropping "zipcode_98003" with p-value 0.9081859817590578
Dropping "zipcode_98188" with p-value 0.780713820714147
Dropping "floors_20" with p-value 0.6356104612914132
Dropping "grade_3" with p-value 0.5856472881210619
Dropping "zipcode_98002" with p-value 0.5406474710298468
Dropping "floors_35" with p-value 0.349191640439133
Dropping "floors_25" with p-value 0.32369742178652816
Dropping "zipcode_98148" with p-value 0.15467038496210414
Dropping "zipcode_98198" with p-value 0.21086855197031598
```

```
In [15]: get_model_data(model, True)
plt.tight_layout()
plt.savefig('figures/final-qq-plot.png')
plt.show()
```

```
R2 = 0.8566083994927416
R2 adjusted = 0.8558099143415273
RMSE (train) = 103086.41637514034
RMSE (test) = 99701.98273002672
number of significant features = 90
```



```
In [16]: model.summary()
```

```
Out[16]: OLS Regression Results
```

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.857
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.856
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1073.
<b>Date:</b>	Sun, 29 Nov 2020	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	10:31:48	<b>Log-Likelihood:</b>	4484.5
<b>No. Observations:</b>	15892	<b>AIC:</b>	-8791.
<b>Df Residuals:</b>	15803	<b>BIC:</b>	-8108.
<b>Df Model:</b>	88		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	-86.9633	7.303	-11.908	0.000	-101.278	-72.649

## Adding Interaction Features

We use the combinations function to create every possible combination of our continuous variables. Looping over them, we create a new interaction variable in our training set, dropping the two

originals. We then make a new model with this train set, and check its R2 and adjusted R2.

```
In [17]: from itertools import combinations
```

```
In [18]: interactions = pd.DataFrame(columns = ['interaction', 'r2', 'r2_adj'])
interactions = interactions.append({'interaction':'baseline', 'r2':0.856608399492})
continuous = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'sqft_basement']
feat_combinations = combinations(continuous, 2)
for i, (a, b) in enumerate(feat_combinations):
    train_interactions = train.copy()
    train_interactions['interaction'] = train_interactions[a] * train_interactions[b]
    train_interactions.drop([a, b], axis=1, inplace=True)

    predictors_interactions = '+'.join(train_interactions.columns[1:])
    formula_interactions = 'price' + '~' + predictors_interactions
    model_interactions = ols(formula=formula_interactions, data=train_interactions)
    train_r2, train_r2_adj = model_interactions.rsquared, model_interactions.rsquared_adj

    interactions = interactions.append({'interaction':f'{a}*{b}', 'r2':train_r2,
```

```
In [19]: interactions.sort_values(by=['r2'], ascending=False).head()
```

Out[19]:

	interaction	r2	r2_adj
0	baseline	0.856608	0.855810
26	lat*long	0.856450	0.855660
4	bedrooms*sqft_basement	0.856086	0.855293
7	bedrooms*yr_since_built	0.856004	0.855212
16	sqft_living*lat	0.855948	0.855155

```
In [20]: interactions.sort_values(by=['r2_adj'], ascending=False).head()
```

Out[20]:

	interaction	r2	r2_adj
0	baseline	0.856608	0.855810
26	lat*long	0.856450	0.855660
4	bedrooms*sqft_basement	0.856086	0.855293
7	bedrooms*yr_since_built	0.856004	0.855212
16	sqft_living*lat	0.855948	0.855155

It looks like every interaction decreases both our R2 and adjusted R2. So, we will not include any interaction variables in our model.

## Cross Validation

We cross validate our data in order to verify that we aren't overfitting to our test set, and that there

aren't outliers in our test set skewing the results. statsmodel does not have its own cross validation function, so we split the total data into parts manually, and loop over the resulting folds and calculate the rmse of each fold individually. The resulting values are near enough to each other that we can call this test a success.

```
In [21]: def kfolds(data, k):
# Force data as pandas DataFrame
data = pd.DataFrame(data)
num_observations = len(data)
fold_size = num_observations//k
leftovers = num_observations%k
folds = []
start_obs = 0
for fold_n in range(k):
    if fold_n < leftovers:
        #Fold Size will be 1 larger to account for leftovers
        fold = data.iloc[start_obs : start_obs+fold_size+1]
        start_obs += 1
    else:
        fold = data.iloc[start_obs : start_obs+fold_size]
        start_obs += fold_size
    folds.append(fold)

return folds

In [22]: test_errs = []
k=5
folds = kfolds(pd.concat([train, test]).sample(frac=1, random_state=100), k)
for n in range(k):
    # Split in train and test for the fold
    fold_train = pd.concat([fold for i, fold in enumerate(folds) if i!=n])
    fold_test = folds[n]
    # Fit a linear regression model
    predictors = '+'.join(fold_train.columns[1:])
    formula = 'price' + '~' + predictors
    fold_model = ols(formula=formula, data=fold_train).fit()
    #Evaluate Test Errors
    y_hat_fold_test = fold_model.predict(fold_test.drop('price', axis=1))
    y_fold_test = fold_test['price']
    test_errs.append(np.sqrt(mean_squared_error(np.exp(y_fold_test), np.exp(y_hat_fold_test))))
print(test_errs)


[102263.4216279941, 103702.93643726473, 105491.1945506598, 100820.10288828705,
102073.61737834036]
```

Our cross validation scores are not far off from each other, so we do not have a problem.

## Residual Plots

Our residuals plot tells us where our errors lie. The distribution seems mostly evenly distributed around the horizontal axis, and not skewed left or right, so we have no problems here. We can save the parameters of our model to a csv to analyze later.

```
In [23]: plt.scatter(model.predict(train.drop('price', axis = 1)), model.resid)
plt.plot(model.predict(train.drop('price', axis = 1)), [0 for i in range(len(train))]
plt.title('Final Model Residuals')
plt.tight_layout()
plt.savefig('figures/final-residuals-plot.png')
plt.show()
```



```
In [24]: model.summary()
```

Out[24]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.857
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.856
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1073.
<b>Date:</b>	Sun, 29 Nov 2020	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	10:32:11	<b>Log-Likelihood:</b>	4484.5
<b>No. Observations:</b>	15892	<b>AIC:</b>	-8791.
<b>Df Residuals:</b>	15803	<b>BIC:</b>	-8108.
<b>Df Model:</b>	88		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	-86.9633	7.303	-11.908	0.000	-101.278	-72.649

```
In [25]: model.params.to_csv('data/final_model.csv', header=False)
```



# Final Model Analysis

Finally, it is time to analyze our final model. We will create a prediction function, examine our coefficients, and draw conclusions.

```
In [1]: import pandas as pd  
import numpy as np
```

```
In [2]: df = pd.read_csv('data/final_model.csv', header = None, names = ['names', 'coeffi  
coef = df['coefficients'].to_dict()
```

## Prediction Function

To create our prediction function, we will start off with the data in the format it is in the dataset we were originally given, `king_house_data.csv`. Since we are using it with our coefficients dictionary, we don't actually need to remove the extra data: we can just only use data that is in both dictionaries. We do, however, need to turn all possible data into numeric data types, log transform the necessary data, and inverse log transform (`np.exp`) the price at the end. Some data will need to be kept as a string in order to create dummy variables from it. Finally, we replace all nan values and turn the data into a dictionary. We can check this on a single row, but we also read in the first 5 lines from `king_house_data.csv` and try those.

```
In [3]: def predict(data):
        data = data.split(',')
        columns = 'id,date,price,bedrooms,bathrooms,sqft_living,sqft_lot,floors,water'
        df = pd.DataFrame(dict(zip(columns, data)), index=[0])
        df['yr_sold'] = df.date.map(lambda x: int(x.split('/')[1]))
        df.drop('date', axis=1, inplace=True)
        for col in df.columns:
            df[col] = pd.to_numeric(df[col])
        df['yr_since_renovation'] = np.where(df['yr_renovated']==0.0, df['yr_sold']-df['yr_built'], df['yr_renovated'])
        df['yr_since_built'] = df['yr_sold'] - df['yr_built']
        categoricals = ['floors', 'condition', 'grade', 'zipcode']
        df = df.astype({col: 'str' for col in categoricals})
        df = pd.get_dummies(df)
        df['renovated'] = df.yr_renovated.map(lambda x: 1 if x>0 else 0)
        continuous = ['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot']
        for col in continuous:
            df[col] = df[col].map(np.log)
        df.replace(np.nan, 0, inplace=True)
        data_dict = df.iloc[0].to_dict()

        prediction = coef['Intercept']
        for key, value in coef.items():
            prediction += value * data_dict.get(key, 0)
        prediction = np.exp(prediction)
        return round(prediction, 2)
```

```
In [4]: data = '7129300520,10/13/2014,221900.0,3,1.0,1180,5650,1.0,,0.0,3,7,1180,0.0,1955'
        predict(data)
        # real price value: 221900.0
```

Out[4]: 239005.82

```
In [5]: # todo: import data from kc_house_data.csv and run the prediction function
        with open('data/kc_house_data.csv') as f:
            f.readline()
            for i in range(5):
                print(predict(f.readline()))
        # the real price values are: 221900.0, 538000.0, 180000.0, 604000.0, and 510000.0
```

239005.82  
576006.35  
239994.68  
556154.46  
471701.59

## Coefficient Analysis



- a 1-year increase in the age of a house increases its sale price by .04%

renovated: 0.062242978234706765

- a house that has been renovated has its sale price increased by 6%

floors\_15: 0.015254015266802584 and floors\_30: -0.0635517536782565

- using a one-floor house as a baseline, a 1.5-floor house has its price increased by 1.5%, while a 3-floor house has its price decreased by 6.4%. Other numbers of floors are approximately equal in price to a 1-floor house.

condition\_2: 0.18927766124862488, condition\_3: 0.3134553949417471, condition\_4: 0.3483349587318392, and condition\_5: 0.3986555695222103

- using a condition of 1 as a baseline, a condition of 2 increases the price by 19%, a condition of 3 increases the price of a house by 31%, a condition of 4 increases the price by 35%, and a condition of 5 increases the price by 40%

Using zipcode 98001 as a baseline, the listed zipcodes increase or decrease the price of a house by 100 times the number listed as a percentage. The unlisted zipcodes are all approximately the same in price. (I'm not planning on listing them all out.)

## Recommendations and Future Work

We attempted to add interaction features to our model, but our results indicated that they only decreased the accuracy of our model. With more time, we could take a deeper look at these and find out why that is the case, and see if other interactions could help our model.

Similarly, adding polynomial features could make our predictions more accurate. Trial-and-error would be needed to determine which features could be changed in this way to improve our model.

Using a mapping library could turn the longitude and latitude into more directly beneficial information, like distance to a school or grocery store. With more time, we could create new features using this information to add to our model.

## Conclusions

Our final model will be useful in predicting sale prices of houses in King county. We can use these predictions to help our clients set the prices for their houses, and find houses that are currently underpriced.